

# CouchDB

## NDBIO40

Jakub Sýkora

# Couch DB

## Introduction

- CouchDB: A Distributed, NoSQL Document Database
- Scalable, Fault-Tolerant, and Highly Available
- Built on Web Standards: JSON, HTTP, and JavaScript
- Master-Master Replication & Eventual Consistency

# CouchDB

## Installation

- Requirements: Docker, Python  $\geq 3$
- Installation steps and checkpoints explained in README.md
- The following steps:
  - Run `docker-compose up -d`
  - Run `init-cluster.sh` (MacOS/Linux) or `init-cluster.ps1` (Windows)
  - Run `pip install -r requirements.txt`
  - Run `python init.py`

# CouchDB

## Installation

- Init-cluster script initialises CouchDB in cluster of 3 nodes
- Init.py script inserts data into CouchDB cluster
- For simplicity in all requests we use URL embedded username, password
- To re-execute installation run first 'docker-compose down —volumes'

# CouchDB

## Domain for the Project

- Libraries with books, authors, lease offers
- Using real data harvested from SPARQL WikiData query
- datagen.py generates json data files

# CouchDB

## Data Model

- Data stored as JSON documents
  - Flexible schema, fields can vary between documents
- Documents organised in databases
- Unique identifiers (`_id`) and revision tracking (`_rev`)
  - Each document has unique `_id`
  - `_rev` tracks changes to the document for conflict resolution

# CouchDB

## Inserting Data

- We want to insert the following documents into database 'library'

```
{  
  "_id": "author_3",  
  "type": "author",  
  "authorLabel": "Emil Cioran",  
  "dob": "1911-04-08T00:00:00Z",  
  "dod": "1995-06-20T00:00:00Z"  
}
```

- Best practise is to give each document 'type' - since there is no schema in CouchDB database can be used to store multiple types of docs (books, library, offers, authors)

# CouchDB

## Inserting Data

- Done via HTTP POST requests as any other database operation
- To see detailed example visit `data_insert.py` script
- In `init.py` we use ``couchdb`` Python lib facilitating database operations



# CouchDB

## Replication and CAP Theorem

- Multi-master replication: each node can accept writes
- Data propagation among nodes is asynchronous
- Network partitioning or concurrent updates can cause inconsistency
- CouchDB aligns with the AP properties of the CAP theorem
  - CouchDB employs eventual consistency
- Not suitable for applications requiring real-time data accuracy
  - e.g., financial transactions

# CouchDB

## Viewing Data

- Built-in view called `_all_docs`
- Custom MapReduce views (complex queries, indexed)
- Mango queries (JSON query definition, ad-hoc, possible indexing)

```
"selector": {  
  "type": "lease_offer",  
  "libraryId": "library_2"  
},  
"fields": ["_id", "_rev", "bookId", "libraryId", "availability", "leaseStart", "leaseEnd"]
```

- Query times (find\_query.py)
  - Unindexed Mango query: 0.039917 seconds
  - Indexed Mango query: 0.007773 seconds
  - Map views query: 0.018620 seconds

# CouchDB

## Introduction to MapReduce Views

- Customizable data querying and aggregation
- JavaScript Map and Reduce functions
  - Map: Process each document and emit key-value pairs
  - Reduce: Aggregate and condense the emitted key-value pairs
- Stored in design documents that exist for each database
- Queried using HTTP, using keys created by MapReduce

# CouchDB

## MapReduce View Example

- Create a view that will represent secondary index over library's country
- PUT into "<couchdb\_url>/<database\_name>/\_design/views":

```
{
  "_id": "_design/views",
  "views": {
    "libraries": {
      "map": "function(doc) { if (doc.type === 'library') { emit(doc.country, doc); } }"
    }
  }
}
```

- Query libraries by country (HTTP GET):
  - <couchdb\_url>/<database\_name>/\_design/views/\_view/libraries?key="Germany"

# CouchDB

## MapReduce Views Implementation Details

- Views are stored as B-trees
  - On first query of the view, the B-tree index is built
  - B-tree keys are emitted keys from Map function
  - Instead of “key” you can use “startKey” and “endKey” to impl. range q.
- Scalable and optimized
  - Incremental updates to avoid reprocessing all documents
  - Parallel processing across distributed nodes

# CouchDB

## MapReduce View Limitations

- Limited Query Types (key, startKey, endKey)
- MapReduce views can be resource-intensive
- Learning Curve
- No JOINS or Aggregations
- Index Storage Overhead
- Compaction Requirements

# CouchDB

## Complex Queries: View Collation and include\_docs=true

- SQL Joins can be partly implemented using view collation or include\_docs
- View collation JOIN to get books written by author
  - Book doc: iterate authors array, emit ([author\_id, ], book)
  - Authors doc: emit([author\_id], doc)
  - Query: range ?startkey=["<author\_id>"]&endkey=["<author\_id>", {}]'
- Data model of domain restricts allow usage of collation, not all JOINS can be achieved

# CouchDB

## Complex Queries: View Collation and include\_docs=true

- Script authors\_books\_join.py measures three implementations of a query
  - MapReduce view using view collation
  - MapReduce view with include\_docs=true
  - MapReduce views to fetch books and authors, join them externally
- Measured on Apple M2 Max
  - View collation (0.008s) is fastest since it uses index directly
  - Include\_docs=true (0.01s) needs additional time to fetch the documents
  - External join (0.01s) only fast because of local network and query size



# CouchDB

## Comparison with SQL constructs

- Aggregates (e.g. COUNT) can be achieved with MapReduce Views
  - authors\_count\_by\_dob.py (shows also GROUP\_PY option on [MM,DD,YYYY] keys by using group\_level param in query)
- JOINS using MapReduce in simple cases or external JOINS
- WHERE: MapReduce function implementation decides what to emit()
- Sorting: Keys ordering in B+ tree (MapReduce View index)